Technical Components Used in RAG-Based AI Systems

Presented by: Senior AI Architect

Program: Externship Review

Made with GAMMA

System Architecture Overview

High-Level RAG Architecture

Decoupling knowledge retrieval from text generation for enhanced accuracy and relevance.

Retrieval + Generation Integration

Seamless orchestration of information retrieval and large language model response synthesis.

Modular System Design

Independent, interchangeable components allowing for flexible development and updates.

Scalable Structure

Designed to handle increasing data volumes and user queries efficiently.

Made with GAMMA

Data Processing Components

Document Loaders

Ingesting data from diverse sources: PDFs, web pages, databases.

Text Extraction

Converting various file formats into plain, structured text.

Data Cleaning

Removing noise, standardizing formats, and correcting errors for quality input.

Chunking Mechanism

Breaking down documents into manageable, semantically coherent segments for retrieval.

Made with GAMMA

Embedding Model

The embedding model is a critical component that transforms text into numerical vector representations. These vectors capture the semantic meaning of the text, allowing for efficient similarity comparisons.

- **Text-to-Vector Conversion**Translates natural language into high-dimensional numerical embeddings.
- **Semantic Representation**Captures the contextual meaning of words and phrases in vector space.
- **Model Selection**Commonly uses advanced models like OpenAI's embeddings or Sentence Transformers.
- **Embedding Dimension**The fixed size of the vector, e.g., 1536 for OpenAI ada-002, defining its feature space.

Made with GAMMA

Vector Database

A specialized database optimized for storing and querying high-dimensional vectors, enabling rapid similarity searches fundamental to RAG systems.

Vector Storage

Efficiently stores millions or billions of vector embeddings.

Similarity Search

Rapidly finds vectors most similar to a query vector.

ANN Indexing

Utilizes Approximate Nearest Neighbor algorithms like HNSW or IVF for speed.

Tools

Popular choices include FAISS, Chroma, Pinecone, and Weaviate.

Made with GAMMA

Retrieval Engine

The retrieval engine is responsible for fetching the most relevant information from the vector database based on the user's query.

- **Semantic Search**Finds documents based on meaning, not just keyword matching.
- **Hybrid Retrieval**Combines vector similarity with traditional methods (e.g., BM25) for robust results.
- **Top-K Tuning**Optimizing the number of top-ranked documents to retrieve for context.
- **Metadata Filtering**Refining search results using associated document attributes.
- **Query Expansion**Augmenting user queries with synonyms or related terms for broader coverage.

Made with GAMMA

Language Model Integration

The LLM component synthesizes the retrieved context and user query into coherent, accurate, and contextually rich responses, minimizing factual errors.

Context Injection

Feeding retrieved relevant information directly into the LLM's prompt.

Prompt Engineering

Crafting effective prompts to guide the LLM's response generation.

LLM Response Generation

Leveraging the LLM to formulate an answer based on the provided context.

Reducing Hallucination

Mitigating the LLM's tendency to generate fabricated information.

Model Used

Integration with state-of-the-art models like GPT series or Google Gemini.

Made with GAMMA

Deployment & Tools

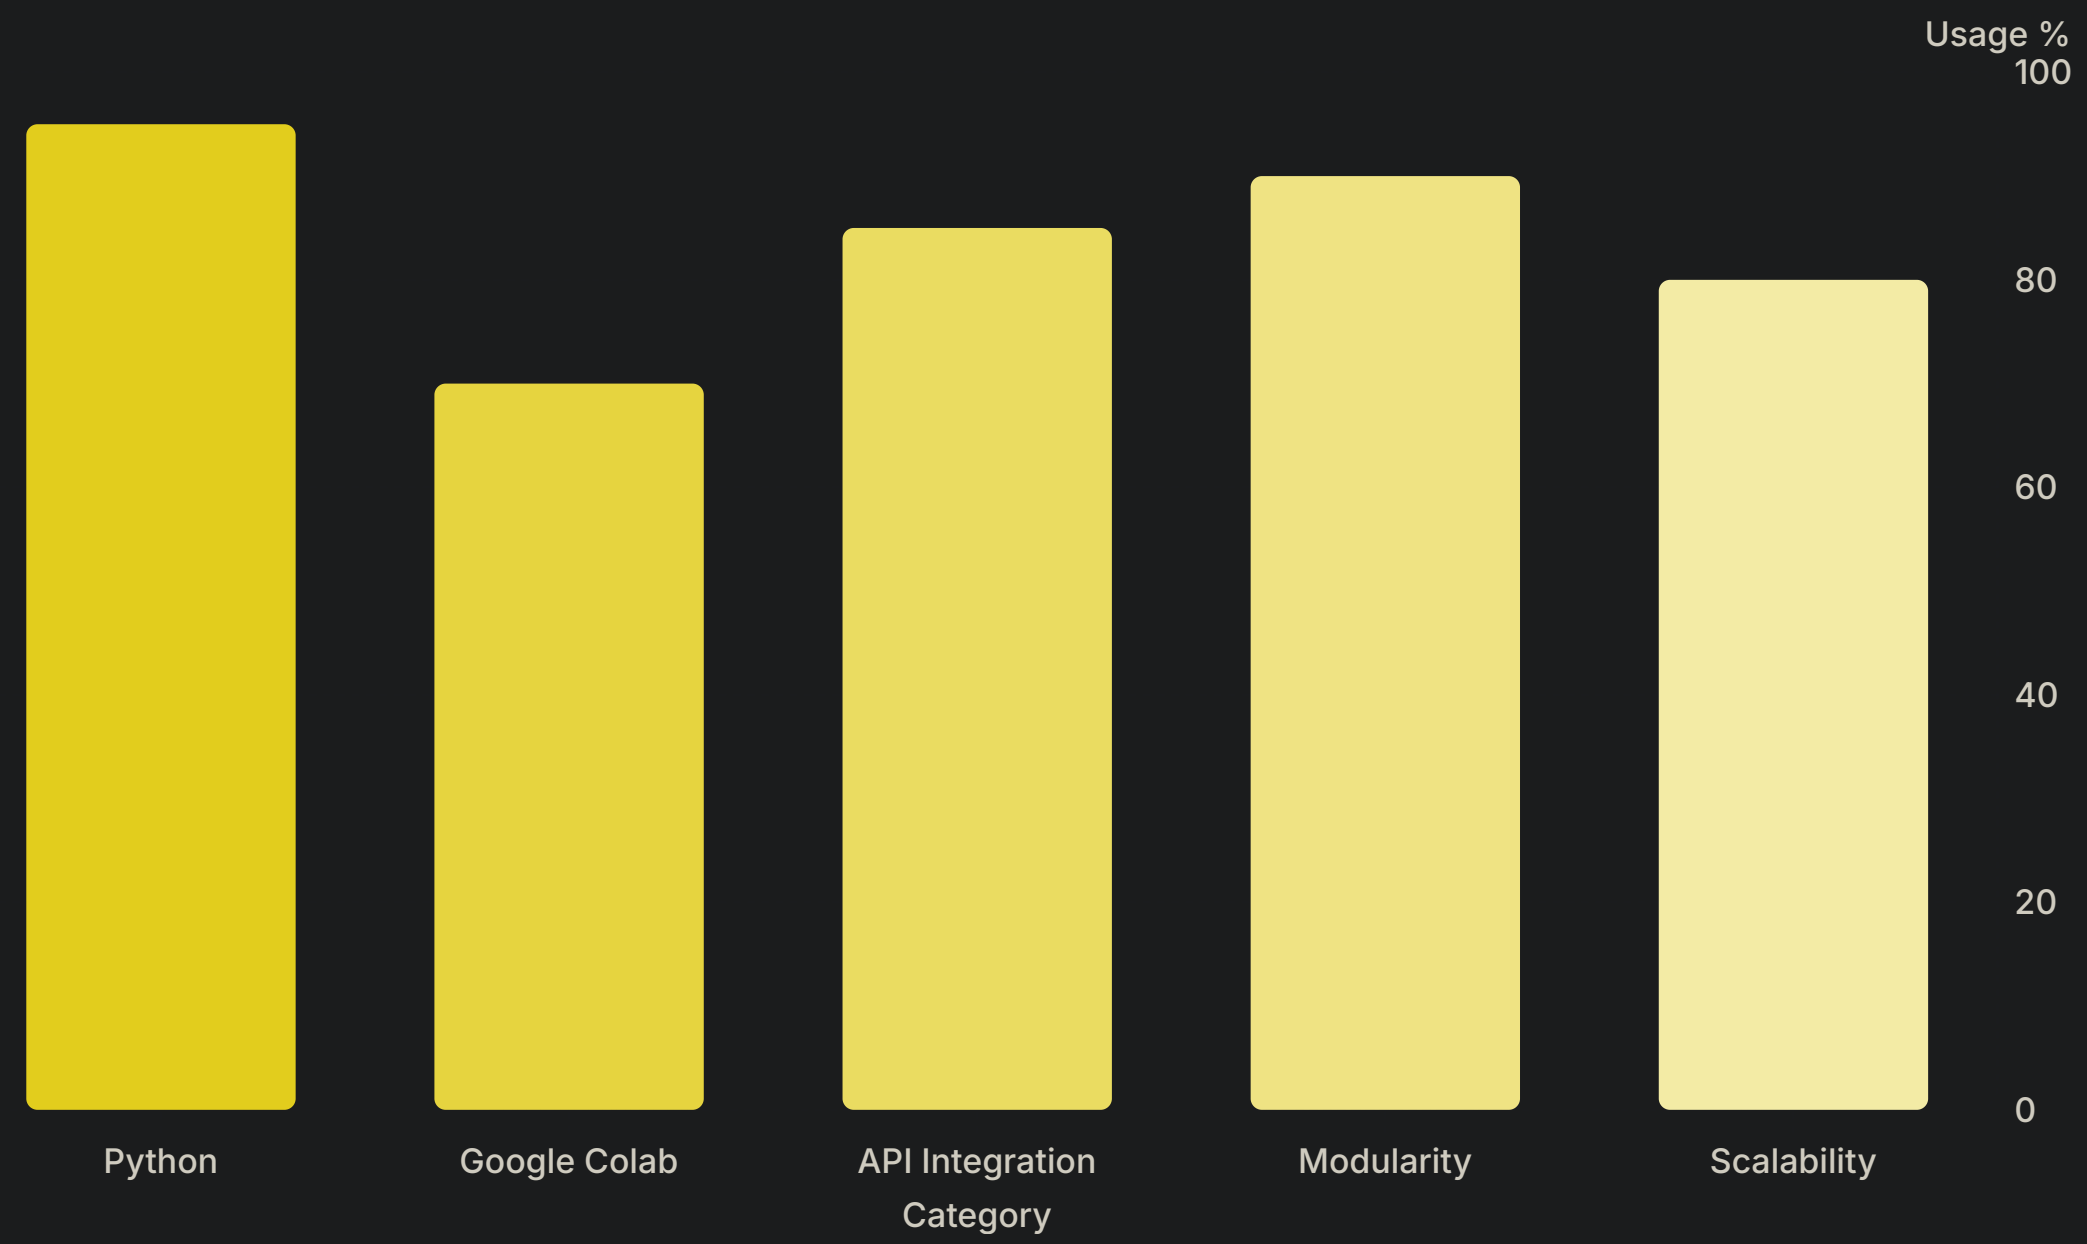

Effective deployment strategies and a robust toolchain are essential for bringing RAG systems from development to production, ensuring reliability and performance.

- **Python Environment**The primary language and ecosystem for RAG development, leveraging extensive libraries.
- **Google Colab**Often used for rapid prototyping, experimentation, and collaborative development.
- **API Integration**Connecting various RAG components and external services via well-defined APIs.
- **Modular Pipeline Structure**Designing pipelines with distinct, independent stages for maintainability.
- **Scalable Deployment Capability**Utilizing cloud platforms and containerization for high availability and elastic scaling.

Made with GAMMA